

HW 4

Q1

- See hw4-nhvasan.zip for more details.
- See hw4-nhvasan.zip for more details.
- See hw4-nhvasan.zip for more details.
- See hw4-nhvasan.zip for more details.

Q2

- See code for implementation details, but here is the model's performance on the Binary and Count Datasets with 200 epochs.

Binary:

```
{0: 2892, 1: 1307, 2: 1307, 3: 187, 4: 268, 5: 198, 6: 226, 7: 154, 8: 142, 9: 116, 10: 114, 11: 110, 12: 102, 13: 104, 14: 85, 15: 81, 16: 74, 17: 77, 18: 66, 19: 68,
20: 59, 21: 62, 22: 56, 23: 58, 24: 52, 25: 49, 26: 45, 27: 40, 28: 38, 29: 35, 30: 44, 31: 42, 32: 49, 33: 51, 34: 57, 35: 46, 36: 44, 37: 25, 38: 27, 39: 22, 40: 23
, 41: 26, 42: 21, 43: 25, 44: 21, 45: 25, 46: 23, 47: 18, 48: 16, 49: 19, 50: 16, 51: 16, 52: 18, 53: 17, 54: 20, 55: 18, 56: 17, 57: 18, 58: 19, 59: 22, 60: 18, 61: 1
6, 62: 15, 63: 18, 64: 19, 65: 21, 66: 15, 67: 12, 68: 13, 69: 12, 70: 17, 71: 15, 72: 10, 73: 12, 74: 14, 75: 11, 76: 11, 77: 13, 78: 13, 79: 12, 80: 12, 81: 9, 82: 1
0, 83: 9, 84: 11, 85: 10, 86: 8, 87: 9, 88: 8, 89: 5, 90: 8, 91: 7, 92: 15, 93: 7, 94: 7, 95: 8, 96: 8, 97: 10, 98: 6, 99: 7, 100: 7, 101: 6, 102: 7, 103: 8, 104: 10,
105: 19, 106: 9, 107: 6, 108: 7, 109: 8, 110: 7, 111: 7, 112: 6, 113: 5, 114: 7, 115: 5, 116: 5, 117: 7, 118: 5, 119: 6, 120: 5, 121: 5, 122: 4, 123: 5, 124: 4, 125: 4
, 126: 4, 127: 4, 128: 4, 129: 5, 130: 3, 131: 5, 132: 3, 133: 5, 134: 4, 135: 4, 136: 8, 137: 4, 138: 4, 139: 4, 140: 4, 141: 4, 142: 4, 143: 4, 144: 4, 145: 5, 146:
3, 147: 3, 148: 3, 149: 2, 150: 3, 151: 5, 152: 8, 153: 10, 154: 14, 155: 4, 156: 3, 157: 2, 158: 2, 159: 2, 160: 3, 161: 2, 162: 3, 163: 3, 164: 2, 165: 3, 166: 3, 16
7: 2, 168: 3, 169: 3, 170: 2, 171: 2, 172: 2, 173: 1, 174: 3, 175: 2, 176: 3, 177: 2, 178: 1, 179: 2, 180: 2, 181: 1, 182: 4, 183: 2, 184: 2, 185: 2, 186: 2, 187: 1, 1
88: 3, 189: 2, 190: 2, 191: 1, 192: 3, 193: 1, 194: 2, 195: 1, 196: 1, 197: 1, 198: 1, 199: 1}
Number of mistakes on the test dataset
32
```

Count:

```
{0: 2892, 1: 1307, 2: 1268, 3: 2736, 4: 1288, 5: 722, 6: 545, 7: 672, 8: 2690, 9: 1263, 10: 428, 11: 681, 12: 700, 13: 2058, 14: 1193, 15: 337, 16: 348, 17: 308, 18: 3
32, 19: 296, 20: 323, 21: 410, 22: 645, 23: 845, 24: 1909, 25: 1166, 26: 349, 27: 371, 28: 412, 29: 458, 30: 751, 31: 742, 32: 1288, 33: 929, 34: 672, 35: 589, 36: 763
, 37: 621, 38: 741, 39: 603, 40: 660, 41: 548, 42: 564, 43: 467, 44: 356, 45: 230, 46: 233, 47: 219, 48: 214, 49: 210, 50: 207, 51: 205, 52: 211, 53: 200, 54: 207, 55:
201, 56: 209, 57: 198, 58: 205, 59: 199, 60: 199, 61: 193, 62: 199, 63: 191, 64: 198, 65: 190, 66: 198, 67: 192, 68: 217, 69: 190, 70: 203, 71: 198, 72: 209, 73: 193,
74: 197, 75: 190, 76: 204, 77: 187, 78: 184, 79: 180, 80: 177, 81: 181, 82: 182, 83: 234, 84: 181, 85: 185, 86: 179, 87: 225, 88: 176, 89: 184, 90: 176, 91: 226, 92:
175, 93: 171, 94: 169, 95: 195, 96: 169, 97: 174, 98: 175, 99: 248, 100: 196, 101: 285, 102: 397, 103: 588, 104: 1421, 105: 1111, 106: 560, 107: 542, 108: 713, 109: 65
2, 110: 968, 111: 741, 112: 709, 113: 522, 114: 409, 115: 308, 116: 221, 117: 179, 118: 175, 119: 174, 120: 165, 121: 164, 122: 171, 123: 156, 124: 172, 125: 155, 126:
166, 127: 153, 128: 162, 129: 155, 130: 178, 131: 155, 132: 142, 133: 153, 134: 153, 135: 172, 136: 151, 137: 144, 138: 151, 139: 152, 140: 173, 141: 158, 142: 145, 1
43: 144, 144: 142, 145: 142, 146: 138, 147: 140, 148: 139, 149: 141, 150: 140, 151: 142, 152: 142, 153: 142, 154: 139, 155: 140, 156: 138, 157: 141, 158: 138, 159: 143
, 160: 140, 161: 138, 162: 133, 163: 138, 164: 137, 165: 143, 166: 136, 167: 186, 168: 141, 169: 132, 170: 142, 171: 141, 172: 140, 173: 139, 174: 132, 175: 135, 176:
129, 177: 133, 178: 119, 179: 131, 180: 123, 181: 126, 182: 135, 183: 118, 184: 127, 185: 124, 186: 125, 187: 121, 188: 129, 189: 122, 190: 115, 191: 117, 192: 115, 19
3: 118, 194: 113, 195: 124, 196: 123, 197: 118, 198: 113, 199: 117}
Number of mistakes on the test dataset
71
```

The model converges slower on the Count dataset, resulting in a higher number of mistakes on the test data compared to the Binary dataset. However, this is to be expected because the Count (non-binary) data holds more information, and since

the number mistakes is on a downward progression on average as the number of epochs increases, this result makes sense.

- b. For my K-Fold Cross Validation implementation, I chose $k=10$ as that is a value that is commonly chosen for k . I first determined the optimal number of epochs for both datasets.

Optimal Epochs

- Binary: The optimal number of epochs based off of number of mistakes is 10,000.

	Epoch	Num_Mistakes
0	1	130.7
1	10	14.6
2	20	10.7
3	50	8.7
4	100	9.7
5	250	9.2
6	500	8.7
7	1000	8.7
8	2500	8.7
9	5000	8.5
10	10000	8.2
	Epoch	Num_Mistakes
10	10000	8.2

- Count: The optimal number of epochs based off of number of mistakes is 2500.

	Epoch	Num_Mistakes
0	1	130.7
1	10	65.8
2	20	67.9
3	50	22.1
4	100	24.7
5	250	16.5
6	500	13.6
7	1000	13.1
8	2500	12.2
9	5000	12.2
10	10000	12.6
	Epoch	Num_Mistakes
8	2500	12.2
9	5000	12.2

I then trained the model using that number of epochs, and reported both the test and training error (number of mistakes) for both datasets.

Binary

- Train: 0
- Test: 32

```
Number of mistakes on the train dataset
0
Number of mistakes on the test dataset
32
```

Count

- Train: 9
- Test: 48

```
Number of mistakes on the train dataset
9
Number of mistakes on the test dataset
48
```

c. See the word lists for the top 15 words with the highest and lowest weights for both the Binary and Count datasets below.

Binary:

- Positive: ['on', 'haven', 'dial', 'might', 'ar', 'the', 'achiev', 'numbertnumb', 'each', 'email', 'without', 'split', 'current', 'gordon', 'like']
- Negative: ['client', 'simpli', 'friend', 'those', 'back', 'specif', 'septemb', 'chat', 'remain', 'should', 'do', 'send', 'link', 'apolog', 'famili']

```
Top 15 Positive Weights: ['on', 'haven', 'dial', 'might', 'ar', 'the', 'achiev', 'numbertnumb', 'each', 'email', 'without', 'split', 'current', 'gordon', 'like']
Top 15 Negative Weights: ['client', 'simpli', 'friend', 'those', 'back', 'specif', 'septemb', 'chat', 'remain', 'should', 'do', 'send', 'link', 'apolog', 'famili']
```

Count:

- Positive: ['haven', 'method', 'that', 'on', 'dial', 'numbertnumb', 'keyboard', 'ascii', 'yourself', 'weekli', 'make', 'promot', 'brent', 'pc', 'each']
- Negative: ['flat', 'few', 'request', 'fact', 'next', 'deathtosпамdeathtosпамdeathtosпам', 'septemb', 'accept', 'do', 'those', 'them', 'simpli', 'back', 'specif', 'apolog']

```
Top 15 Positive Weights: ['haven', 'method', 'that', 'on', 'dial', 'numbertnumb', 'keyboard', 'ascii', 'yourself', 'weekli', 'make', 'promot', 'brent', 'pc', 'each']
Top 15 Negative Weights: ['flat', 'few', 'request', 'fact', 'next', 'deathtosпамdeathtosпамdeathtosпам', 'septemb', 'accept', 'do', 'those', 'them', 'simpli', 'back', 'specif', 'apolog']
```

Q3

```
(CS334) nikhitas-mbp-3:hw4-nhvasan niki$ python q3.py
Logit Number of Mistakes [Binary Data] 29
Logit Number of Mistakes [Count Data] 39
Naive Bayes Number of Mistakes [Binary Data] 52
Naive Bayes Number of Mistakes [Count Data] 55
```

- The Naive Bayes model performs slightly better on the Binary dataset than the count, with 52 mistakes on the test Binary dataset versus 55 on the test Count dataset.
- The Logistic Regression model performs better on the Binary dataset as well, with 29 mistakes on the test Binary dataset versus 39 on the test Count dataset.